

Placement Empowerment Program
Cloud Computing and DevOps Centre

IMPLEMENT AUTO-SCALING IN THE CLOUD

(Set up an auto-scaling group for your cloud VMs to handle variable workloads)

NAME: Swathika K

DEPARTMENT: AML

INTRODUCTION:

Auto Scaling is an AWS services that ensure high availability, fault tolerance, and cost efficiency in cloud applications.

Auto Scaling (ASG - Auto Scaling Group):

Auto Scaling automatically adjusts the number of EC2 instances in response to demand.

What Auto Scaling Does:

- **Increases Instances (Scale Out):** When traffic or CPU usage increases, Auto Scaling launches new EC2 instances.
- **Decreases Instances (Scale In):** When traffic decreases, Auto Scaling terminates extra instances to reduce costs.
- **Ensures Availability:** If an instance fails (due to crash, hardware failure, etc.), Auto Scaling replaces it.
- **Optimizes Costs:** Ensures you only use the required number of instances at any given time.
- **Elastic Load Balancer** distributes incoming traffic across multiple EC2 instances to improve performance and availability.

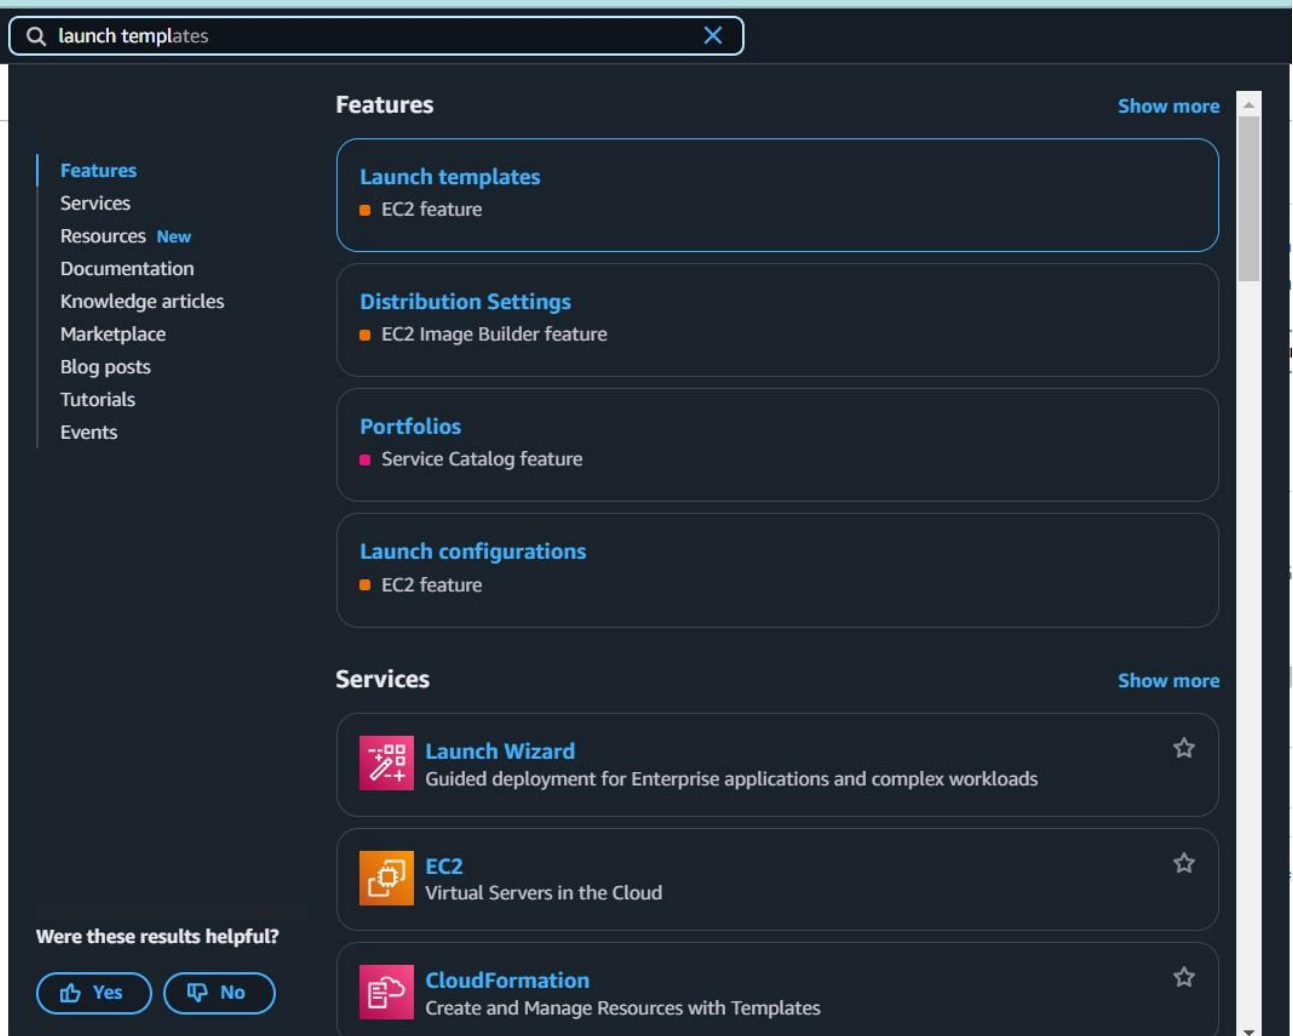
IMPORTANCE:

- Auto Scaling Group monitors load (CPU, traffic) and adds/removes instances as needed.
- If traffic increases, ASG launches more instances, and ELB distributes traffic evenly.
- If traffic decreases, ASG removes extra instances, reducing costs.
- **Result:** High availability, better performance, and optimized cloud costs!

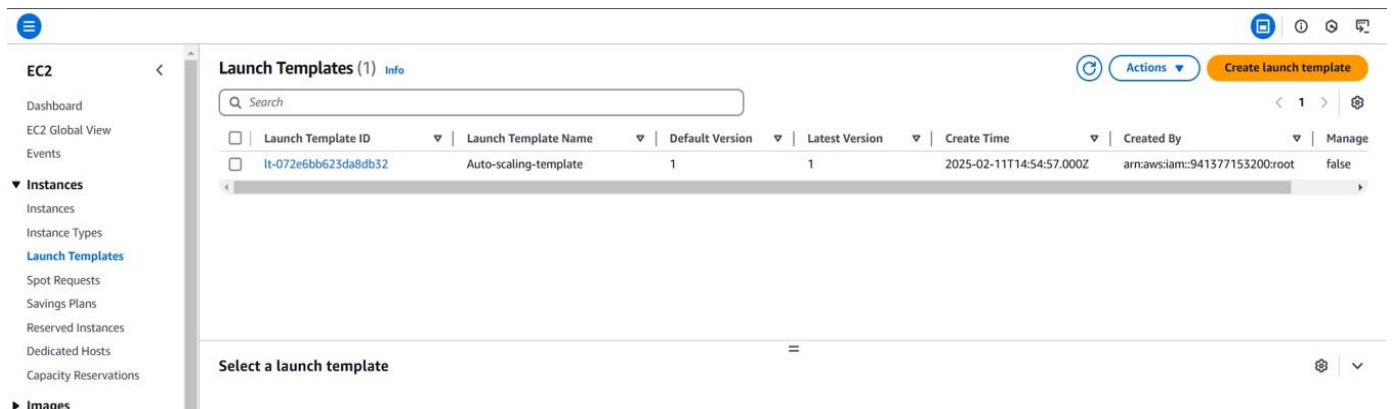
STEP BY STEP OVERVIEW:

Step 1: LAUNCH TEMPLATES

- Login into your AWS console and search for launch templates.



- Navigate to it and click on launch template.
- Specify the template name, AMI image (Amazon Linux), Instance type (t2.micro), Key pair.
- Mainly, configure the security groups and add the security group role(such as: HTTP, ssh)
- Now launch the template.



Step 2: AUTO SCALING GROUPS

- Go to your EC2 dashboard, on the left sidebar you will find 'Auto Scaling groups'. Click on create.
- Give the name to the ASG and select the template you have just created.

The screenshot shows the 'Create Auto Scaling group' wizard in the AWS Management Console, specifically Step 2: Choose instance launch options. The breadcrumb navigation at the top reads 'EC2 > Auto Scaling groups > Create Auto Scaling group'. A vertical progress bar on the left lists seven steps: Step 2 (selected), Step 3 (optional), Step 4 (optional), Step 5 (optional), Step 6 (optional), Step 7, and Review. The main content area is titled 'Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group.' It features a 'Name' section with a text input field containing 'My-ASG' and a note that the name must be unique and under 255 characters. Below this is the 'Launch template' section, which includes an informational box about the EC2 console's support for launch templates as of May 31, 2023. The 'Launch template' dropdown is set to 'Auto-scaling-template', and the 'Version' dropdown is set to 'Default (1)'. A table at the bottom displays details for the selected launch template: 'AMI ID' is 'ami-0c50b6f7dc3701ddd', 'Launch template' is 'Auto-scaling-template' with ID 'lt-072e6bb623da8db32', 'Instance type' is 't2.micro', and 'Request Spot Instances' is 'No'.

Choose your VPC or the default VPC and select two subnets of two different availability zone.

The screenshot shows the 'Create Auto Scaling group' wizard in the AWS Management Console, specifically Step 7: Review. The breadcrumb navigation at the top reads 'EC2 > Auto Scaling groups > Create Auto Scaling group'. The vertical progress bar on the left shows Step 7 as the current step, with 'Add tags' as the sub-step. The main content area is titled 'Network' and provides information about using multiple Availability Zones. It includes a 'VPC' section with a dropdown menu showing 'vpc-02a6d6d355b6ee357' (Default). The 'Availability Zones and subnets' section shows two selected subnets: 'ap-south-1a | subnet-013d52f393abc4060' and 'ap-south-1b | subnet-0ee18172a6bd04005'. The 'Availability Zone distribution' section has two radio button options: 'Balanced best effort' (selected) and 'Balanced only'. At the bottom right, there are four buttons: 'Cancel', 'Skip to review', 'Previous', and 'Next'.

Leave the next setting as default. Now, review and click on create Auto Scaling group.

Auto Scaling groups (1) [Info](#)[Launch configurations](#)[Launch templates](#)[Actions](#)[Create Auto Scaling group](#)

< 1 > ⚙️

<input type="checkbox"/>	Name	Launch template/configuration	Instances	Status	Desired capacity	Min	Max	Availability Zones
<input type="checkbox"/>	My-ASG	Auto-scaling-template Version Default	0	⏸ Updating capacity...	1	1	1	ap-south-1b, ap-south-1a

My-ASG

My-ASG Capacity overview [Edit](#)[arn:aws:autoscaling:ap-south-1:941377153200:autoScalingGroup:0862e805-b873-4a9e-b376-373cfd1888bc:autoScalingGroupName/My-ASG](#)

Desired capacity 1	Scaling limits (Min - Max) 1 - 1	Desired capacity type Units (number of instances)	Status -
-----------------------	-------------------------------------	--	-------------

Date created
Tue Feb 11 2025 20:36:24 GMT+0530 (India Standard Time)

[Details](#) | [Integrations - new](#) | [Automatic scaling](#) | [Instance management](#) | [Instance refresh](#) | [Activity](#) | [Monitoring](#)
Launch template [Edit](#)

Launch template
[lt-072e6bb623da8db32](#)
Auto-scaling-template

Version
Default

Description
-

[View details in the launch template console](#)

AMI ID
[ami-0c50b6f7dc3701ddd](#)

Security groups
-

Storage (volumes)
-

Instance type
t2.micro

Security group IDs
[sg-02a787c226b5e51b6](#)

Key pair name
sample

Owner
[arn:aws:iam::941377153200:root](#)

Create time
Tue Feb 11 2025 20:24:57 GMT+0530 (India Standard Time)

Request Spot Instances
No

Step 3: TESTING AUTO SCALING

**Important Note: Do Not Perform This Test If You Want to Avoid Costs:**

1. Launching and running additional EC2 instances will incur charges beyond the AWS Free Tier.
2. Simulating high CPU usage and triggering scaling may increase costs temporarily due to additional resource allocation.

TESTING:

1. Simulate High CPU Usage on an EC2 Instance: Connect to one of your EC2 instances in the Auto Scaling Group using SSH.
2. Run a command to create artificial CPU load. For example: `sudo yum install -y stress, stress--cpu 2--timeout 300`.

This command will utilize 2 CPU cores for 5 minutes, simulating high CPU usage.

MONITOR SCALING ACTIVITIES:

Navigate to the AWS Management Console > EC2 Dashboard > Auto Scaling Groups.

Select your Auto Scaling Group and go to the Activity History tab.

Check if a new instance is being launched based on your scaling policy (e.g., CPU utilization exceeding 50%).

TERMINATE THE STRESS TEST:

Once testing is done, stop the CPU load by pressing Ctrl+C in the terminal or by terminating the stress process.

VERIFY SCALING DOWN:

After the CPU usage drops, monitor the Auto Scaling Group again to confirm that unnecessary instances are terminated, returning to the desired capacity.

CONCLUSION:

By completing this PoC, we will be able to:

- Implement Auto Scaling group and launch the Template.
- Attain dynamic scaling and monitoring.
- Most importantly, we should be aware of Costs.